

Scheduling

Fabrice Muller

Polytech Nice Sophia


Fabrice.Muller@univ-cotedazur.fr



Introduction

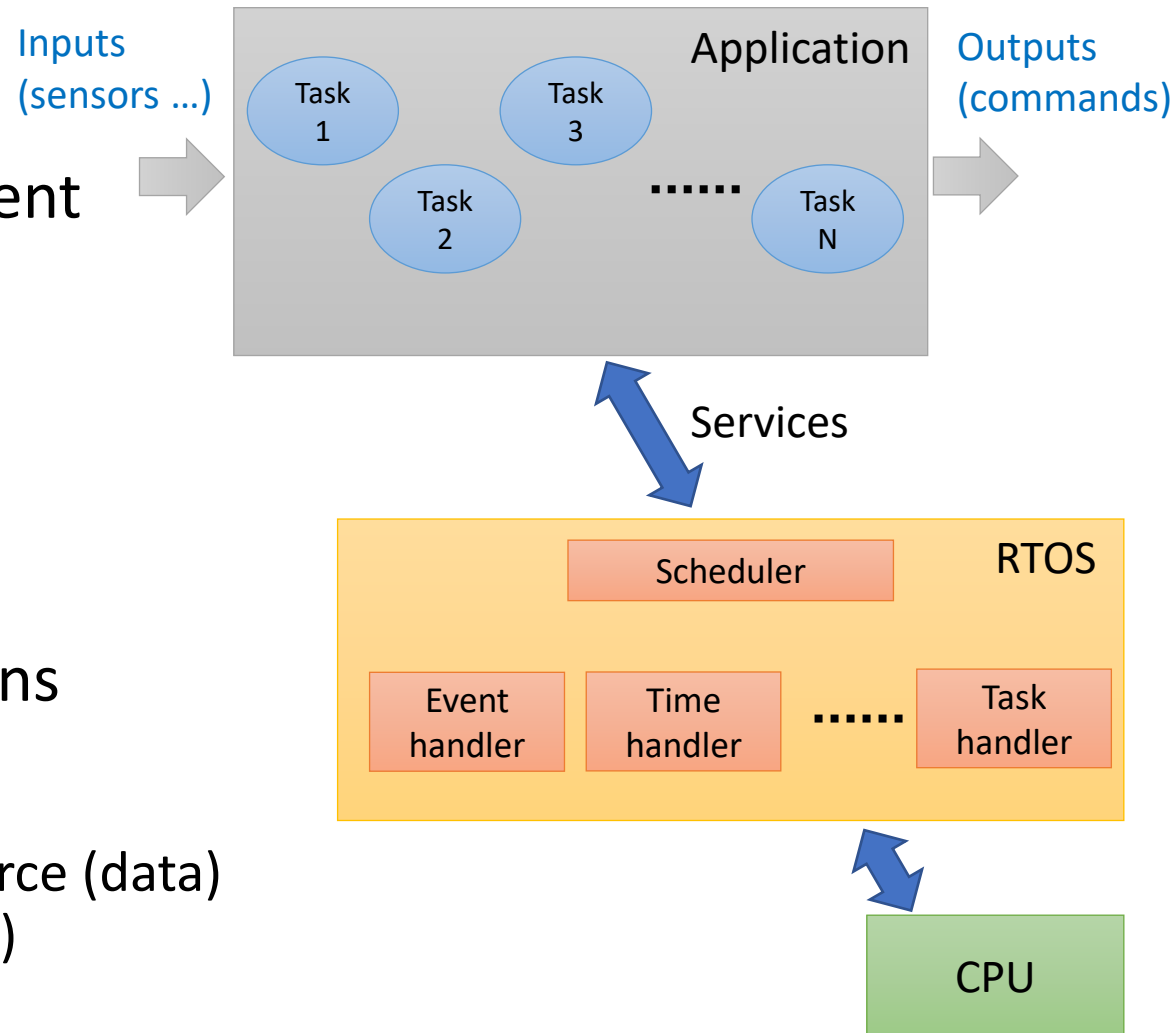
Scheduling without a RTOS

- Simple solution with standalone application
- Use of datas from polling or interrupt solutions
- Libraries of deterministic functions
- Scheduling of tasks by hand





Scheduling with a RTOS

- Need of services
- Input/Output management
 - Input/output handler
 - Interrupt handler
- Task scheduling
 - Organization in tasks
 - Scheduling policy
 - Time handler
- Inter task communications
 - Synchronization (event)
 - Communication (data)
 - Access to a shared resource (data)
 - Time (counter, watchdog)



Inter task communications

- Synchronization
 - Event
 - Semaphore (Boolean)
 - Rendezvous
- Communication
 - Mailbox
 - Semaphore with counter
- Shared critical resources
 - Semaphore with mutual exclusion (Mutex)



Quality of Service (QoS)

- Subjective notion
- Depend on systems and services
 - Response time (interactive or critical applications)
 - Bit rate (video broadcast)
 - Availability (access to shared services)
 - Packet loss rate (voice or video perception) →
 - Signal-to-noise ratio (communication)
 - ...
- Example : For application layer services for telephony and streaming video, QoS is the acceptable cumulative effect on subscriber satisfaction of all imperfections affecting the service.
- **High QoS is often confused with a high level of performance** (high bit rate, low latency and low bit error rate ...)



Criteria for real-time computing

- Hard real time ($\text{QoS} = 100\%$)
 - Missing a deadline = total system failure
 - To be predictable, deterministic and reliable
 - Use of mathematical techniques
- Firm real time ($X\% \leq \text{QoS} \leq 100\%$)
 - Infrequent deadline misses = tolerable ($X\%$)
 - May degrade the system's QoS
 - The usefulness of a result is wrong after its deadline
 - Minimize the probability of missing a deadline several times consecutively
- Soft real time ($\text{QoS} \leq 100\%$)
 - The usefulness of a result degrades the system after its deadline but maybe acceptable

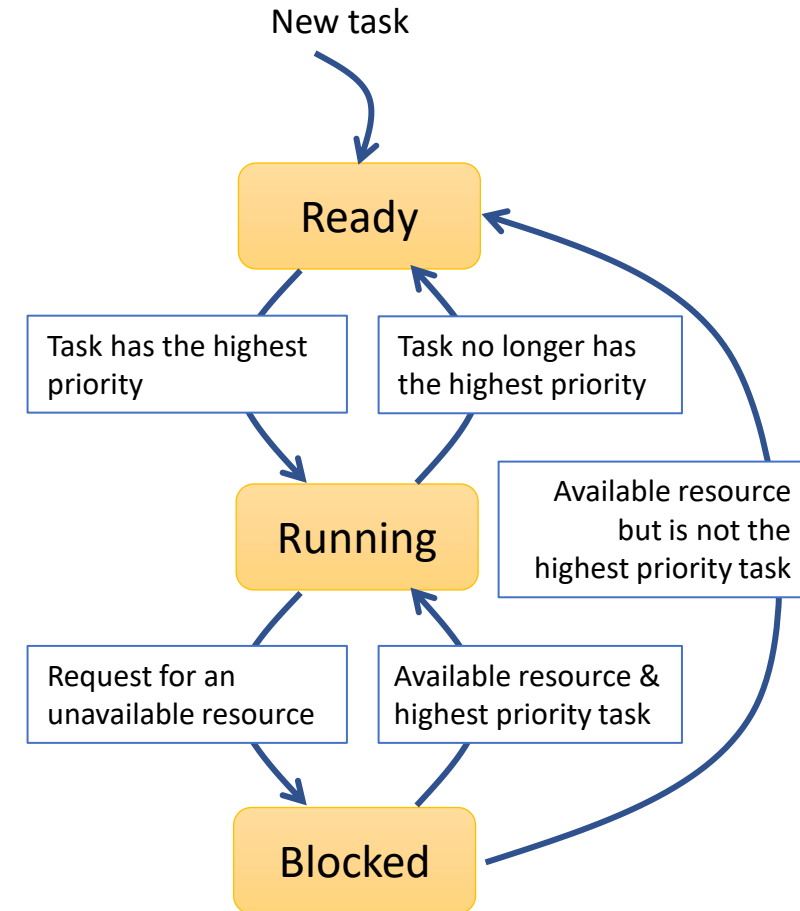
Hard real time - Definitions

- Predictability
 - The performance of the application must be defined in all possible cases to ensure respect for time constraints
 - Use of the worst case
- Determinism
 - No uncertainty about the behavior of the system
 - Behavior is always the same for a given context
- Reliability
 - Ability of a system to achieve and maintain functionality under normal conditions of use
 - Respect of real-time constraints
- Fault Tolerance
 - Fault-tolerance is the ability of a system to maintain its functionality even in the presence of faults
 - Like Reliability even if certain failures have appeared

What is a Task ?

Typical state of a task

- Ready state
 - The task is ready to run but cannot because a higher priority task is executing
- Running state
 - The task is the highest priority task and is running
- Blocked state
 - The task has requested a resource that is not available
 - The task has requested to wait until some event occurs
 - The task has delayed itself for some duration
- Some kernels (VxWorks, FreeRTOS ...) define more granular states such as suspended, pended, delayed ...



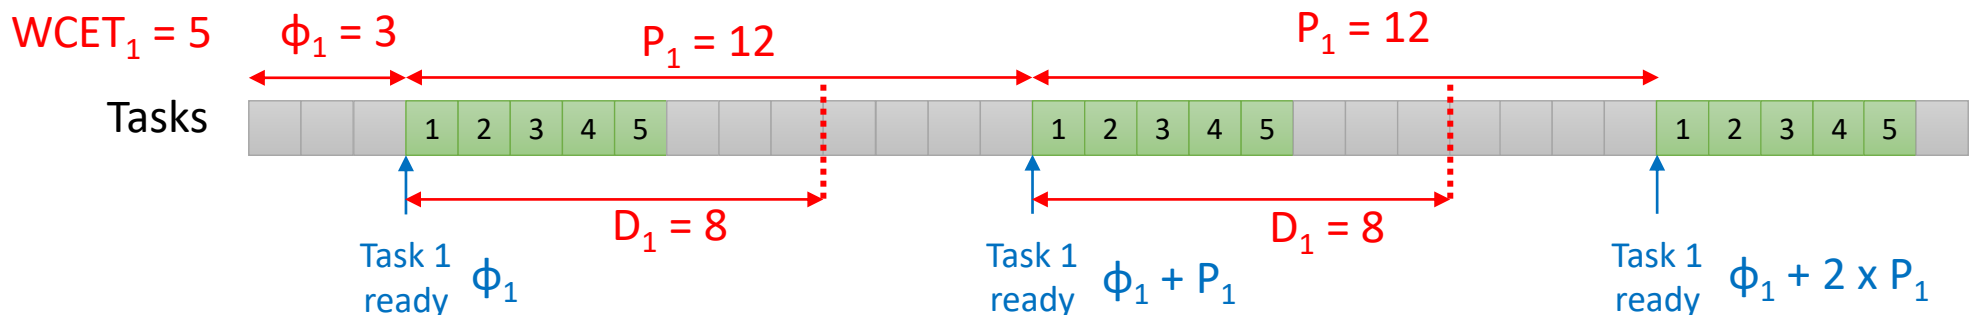
Priority of tasks

- The priority level allows you to define who has access to
 - CPU
 - Shared resources
- The kernel (or tick OS) has the highest priority
- Idle task has the lowest priority (normally zero)
- FreeRTOS Example
 - Each task is assigned a priority from 0 to `configMAX_PRIORITIES-1`
 - Low priority numbers denote low priority tasks

Periodic Tasks

- Task T_i repeats after a certain fixed time interval (period)
- Characteristics of T_i
 - ϕ_i : phase (from 0s till the first execution of the task i)
 - P_i : Period
 - $WCET_i$: Worst Case Execution Time
 - D_i : relative Deadline

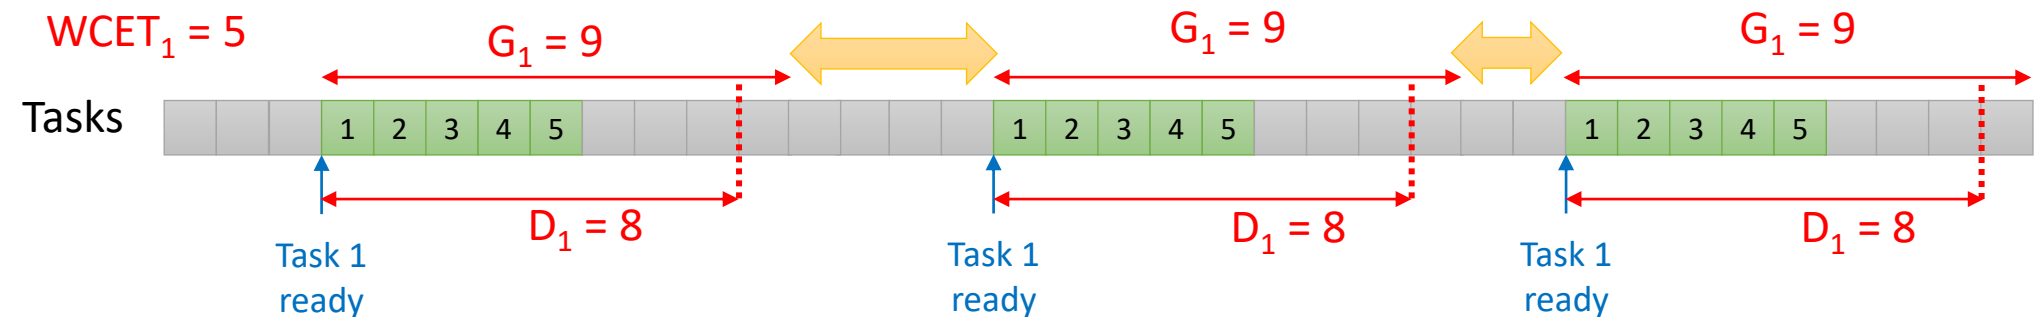
Example of periodic task T_3



Aperiodic / Sporadic Tasks

- Sporadic task
 - Task T_i recurs at random instants with a minimum separation between 2 consecutive instances of task T_i
 - Characteristics of T_i
 - G_i : next instance cannot occur before G_i
 - $WCET_i$: worst case execution time
 - D_i : relative deadline
- Aperiodic task
 - Task T_i can arise at random instants ($G_i=0s$)
 - Two or more instances of a task T_i might occur at the same time
 - Might lead to a few deadline misses (used for firm/soft real-time)

Example of sporadic task T_3



Hyper period - set of periodic tasks

- The execution of any set of periodic tasks is cyclic
- Hyper period of N tasks = $\text{LCM}(\text{Period}_1, \dots, \text{Period}_N)$
- The Least Common Multiple of periods of all the tasks
- Example

$T_1 : P_1 = 29, \text{WCET}_1 = 7$

$T_2 : P_2 = 5, \text{WCET}_2 = 1$

$T_3 : P_3 = 10, \text{WCET}_3 = 2$



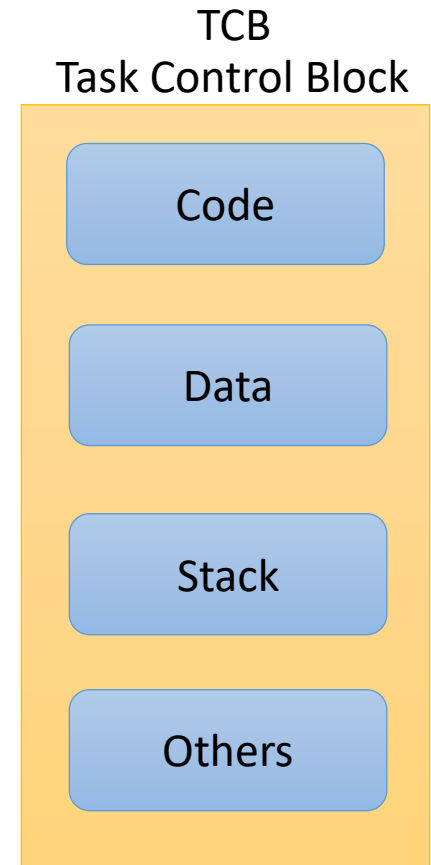
Hyper period = $\text{LCM}(29, 5, 10) = 290$

2 values : $\text{LCM}(a, b) = A \cdot B / \text{GCF}(a, b)$ (using the Greatest Common Factor)

3 values : $\text{LCM}(a, b, c) = \text{LCM}(\text{LCM}(a, b), c) = \text{LCM}(a, \text{LCM}(b, c))$

Task descriptor / Task Control Block

- **Data structure describing a task = TCB** (Task Control Block)
- Scheduler uses TCB for the management of multitask environment
- Code
 - Instructions, constants
- Data
 - Shared with the tasks of the same process
- Stack
 - Contains temporary information
 - Local variables, context, registers
 - Program counter of subroutine
- Others
 - Identifier
 - State
 - Task priority
 - Expected events (rendezvous)
 - ...

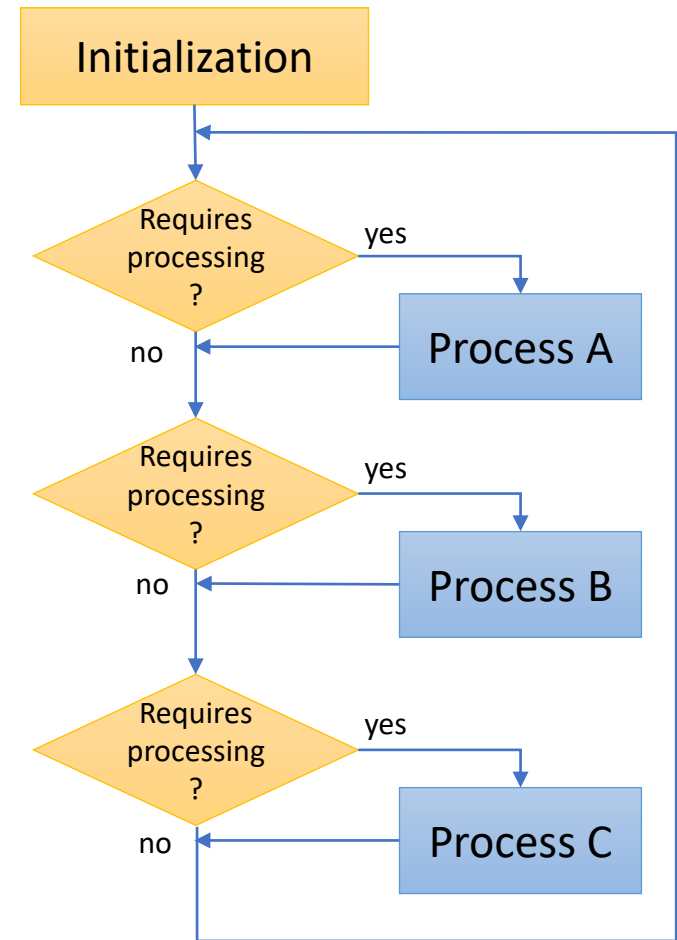


How to manage tasks ?

The task management is done through a scheduler which defines the order of the tasks with their priority.

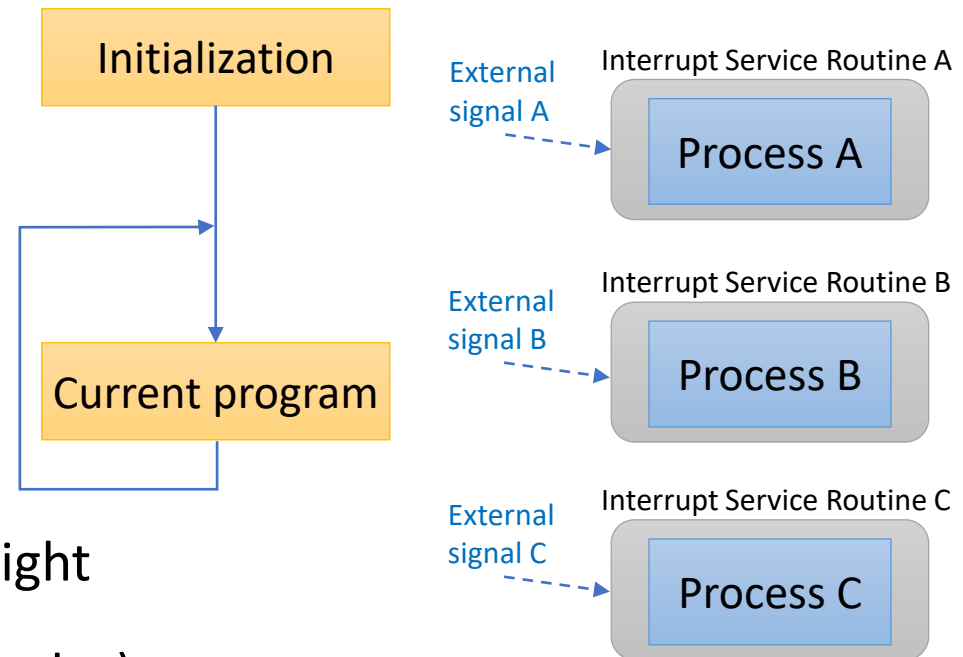
Polling mode

- Polling mode repeatedly checks whether a device needs servicing
- Polling cycle is the time in which each element is monitored once
- Disadvantages
 - if there are too many devices to check, you risk missing a state change of devices.
 - Increases CPU rate (wastes lots of CPU cycles)



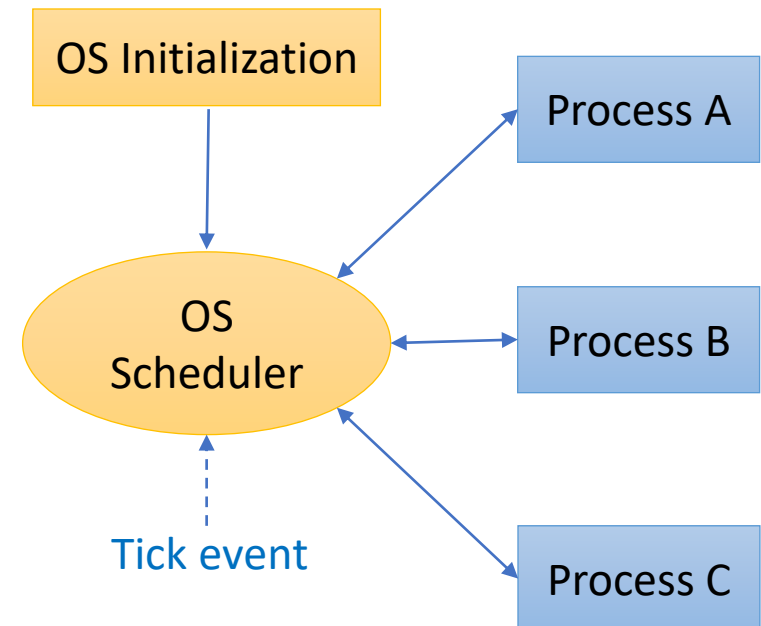
Interrupt mode

- An interrupt is a signal from a device or from a program
 - Stops the current program
 - Saves registers
 - Runs an interrupt routine
 - Restores registers
- For each interrupt
 - Priority assignment
 - Masking to disable interrupt
- Advantages
 - Each process is triggered at the right time
 - Decreases CPU rate (saves CPU cycles)



RTOS mode

- Scheduler
 - Manages process/tasks
 - Can be seen as an interrupt routine
 - Manages context switch (save/restore context in TCB)
 - The highest priority
- Tick event
 - Wake up the scheduler periodically
 - Comes from a timer
 - Tick period is configurable (FreeRTOS = 10 ms by default)




Scheduler

Some definitions

• Tasks

- T • Independent tasks share only the processor
- Dependent tasks
 - T • share other resources
 - T • Have precedence constraints (waiting for data of other tasks)



• Scheduling


- Offline scheduling is pre-calculated before execution
- Online scheduling decides dynamically of the execution of tasks

• Cooperative/Preemptive

- A cooperative scheduler cannot interrupt a task
- A preemptive scheduler can interrupt a task (and save this context) to execute a higher priority task (restore this context)


Goals of scheduler

- Scheduler manages process/tasks
- Scheduler handles
 - the selection of a process/tasks for the processor based on a scheduling algorithm
 - the removal of a process/tasks from the processor
- Ensure properties
 - Equity (all processes must run)
 - Load balancing (multi-core context)
 - ...
- Objectives are often opposed
 - Minimize response time but ...
 - Maximize CPU ratio



Types of Schedulers

- Types of Processes
 - I/O-bound process** - it spends his time doing I/O operations; a lot of little CPU bursts
 - CPU-bound process** – it spends its time doing calculation; small number of very large CPU bursts
- Long-term scheduler / job scheduler
 - Select in the storage pool the processes must be add in the ready queue
 - It must select a careful mixture of I/O bound and CPU bound processes to yield optimum system throughput
 - Not invoked very often (seconds, minutes)
- Short-term scheduler / CPU scheduler
 - Select in the ready queue which process/task should be implemented soon and reserve the CPU
 - Invoked very frequently (in milliseconds), quick response time
- Medium Term Scheduler
 - Remove the processes from the main memory
 - In-charge of handling the swapped out (roll out)-processes
 - Suspended process is moved from ready queue to the storage pool



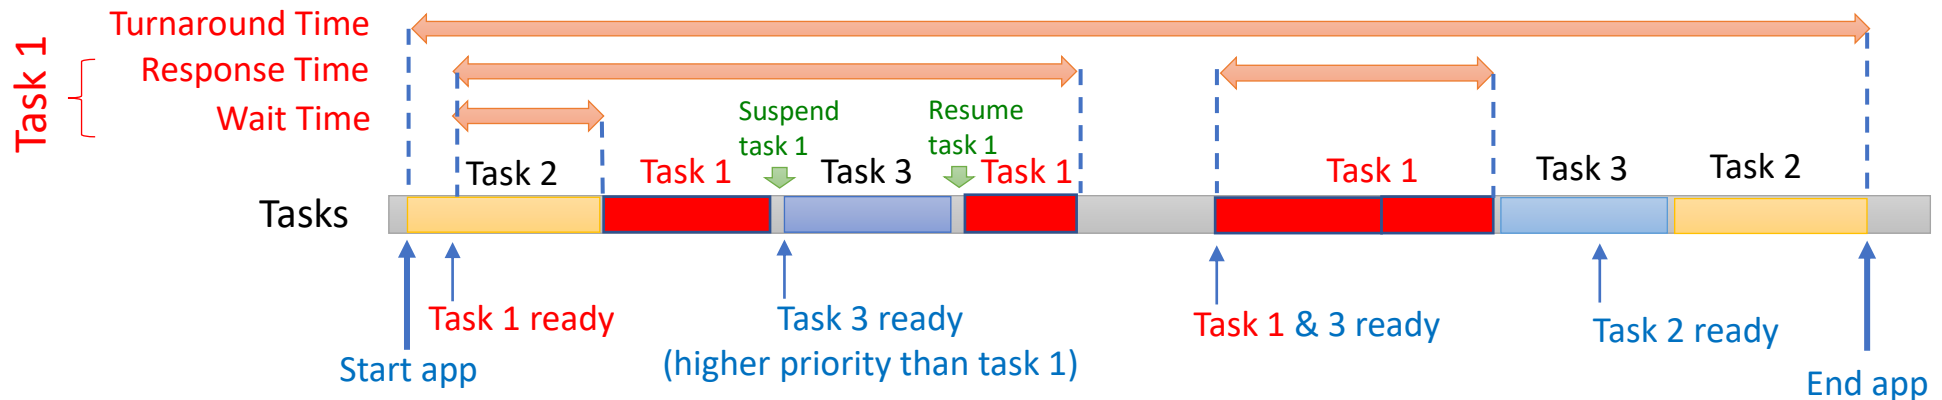
Scheduling Policy

Objectives

- Improving efficiency of the system
 - Reducing delays and wait times, response times to the system
 - Managing CPU resources better
- Fairness
 - Scheduler should not give unfair advantage to a process/task
 - Important to balance long-running tasks and ensure that the lighter tasks can be run quickly
- Heuristic algorithms for scheduling
 - Trade-off between decision time and optimality
 - Heuristic does not always reach optimality
 - Multi-criteria algorithms
- Lot of scheduling policies exist !

Criteria for Scheduling

- Reduce the strain on the **CPU Utilization**
 - Manage the percentage of time the CPU is busy
- Optimize the **Throughput**
 - Increase the number of processes completed in a given time frame
- Reduce the **(Average) Wait Time**
 - Waiting time is amount of time a process has been waiting in the ready queue
- Reduce the **Response Time**
 - The response time of a task/thread is defined as the time elapsed between the dispatch (time when task is ready to execute) to the time when it finishes its job (one dispatch)
- Respect the **Turnaround Time**
 - Total time a process takes to run, from start to finish (includes all waiting time)




Main Execution Times of a task

- Depend on
 - Target architecture (processor, memory accesses, ...)
 - Algorithms (conditions, loops)
- **Worst-Case Execution-Times (WCET)**
 - Maximum length of time the task could take to execute on a specific hardware platform
 - Assess resource needs for real-time systems
 - Ensure meeting deadlines
 - Perform schedulability analysis
- **Best-Case Execution-Times (BCET)**
 - Assess code quality
 - Assess resource needs for non/soft real-time systems
 - Benchmark hardware
- **Average-Case Execution-Times (ACET)**
 - Assess behavior of the real-time systems from simulations

Algorithm

val : range 0 to 255



```

void f1() {
    int val = read(PortA);
    if (val = 0) {
        write(PortB, -1);
    }
    else {
        for (int i=0; i<val, i++)
            write(PortB, var[i]);
    }
}
    
```

Non-Preemptive Scheduling

- Task runs on CPU till it gets terminated or it reaches a blocked state
- Scheduler does not interrupt a task in running state in middle of the execution
- Cooperative scheduler
- Scheduling: Shortest Job First (SJF), Priority


Example of Priority scheduling

Priority of Task 1 (P=10) > Task 2 (P=5) > Task 3 (P=2)

Execution Time

Task 1	1	2	3		
Task 2	1	2	3	4	5
Task 3	1	2	3	4	5

Select Task 1 more
priority than Task 2



First Come First Serve (FCFS)


Non-Preemptive version

- First Come First Serve is just like FIFO (First in First out) Queue
- **Selects task from the head of the queue and new task enters through the tail of the queue**
- Average Waiting Time (AWT)
 - Crucial parameter to judge its performance
 - *Lower the Average Waiting Time, better the scheduling algorithm*
- Disadvantages
 - Non-Preemptive algorithm which means the process priority doesn't matter
 - Not optimal Average Waiting Time
 - Resources utilization in parallel is not possible

Task	Execution Time (ms)
T1	21
T2	3
T3	6
T4	2

$$AWT = (0 + 21 + 24 + 30) / 4 = 18,75 \text{ ms}$$

Tasks in the queue at 0s



Waiting Time

WT(T1) = 0 ms
 WT(T2) = 21 ms
 WT(T3) = 24 ms
 WT(T4) = 30 ms

Shortest Job First (SJF)

Non-Preemptive version

- **Select the ready task with the smallest execution time to execute next**
- Also known as Shortest job next (SJN)
- Greedy Algorithm
- SJF has the advantage of having minimum average waiting time among all scheduling algorithms
- May cause starvation if shorter processes keep coming. This problem can be solved using the concept of aging
- Practically infeasible as Operating System may not know execution time
 - Predict online execution time from previous executions
 - Estimate off-line the execution time for a job for a specific architecture

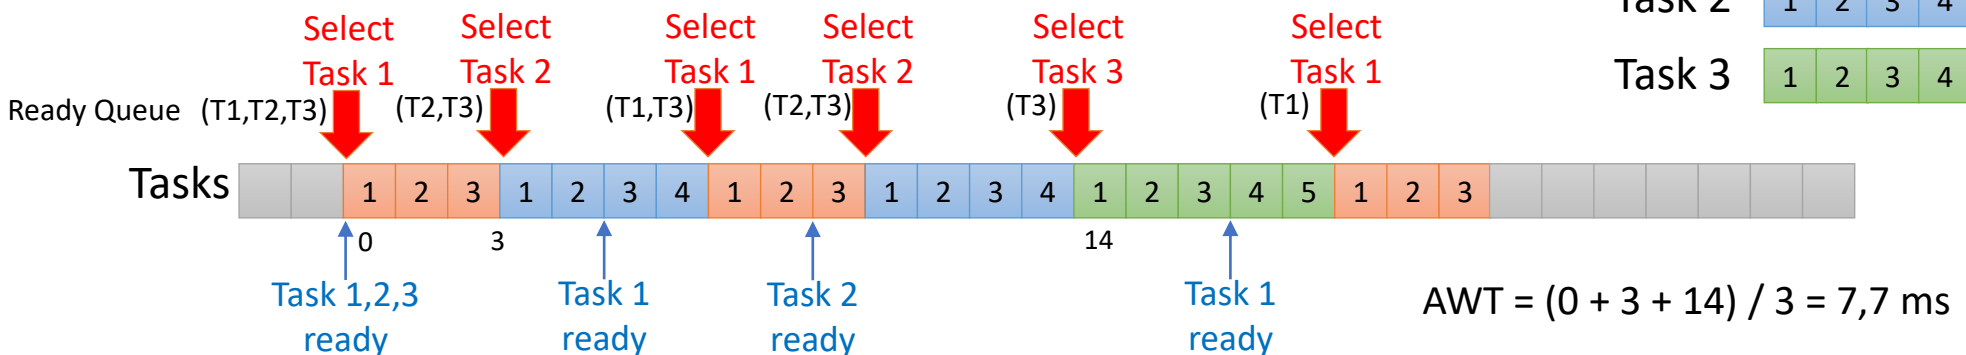
Execution Time

Task 1 1 2 3

Task 2 1 2 3 4

Task 3 1 2 3 4 5

Priority of Task 1 > Task 2 > Task 3
(according to execution time)



Rate-Monotonic scheduling Schedulability

- Priority assignment algorithm
- Properties of tasks
 - No resource sharing = independent tasks
 - Deadlines are exactly equal to periods, $D_i = P_i$
 - Static priorities: **tasks with shorter periods/deadlines are given higher priorities**
 - Task with the highest static priority that is runnable immediately preempts all other tasks
 - Context switch times and other thread operations are free and have no impact on the model
- Schedulability
 - Liu & Layland (1973) proved a feasible schedule that will always meet deadlines exists if the CPU utilization is below a specific bound

$$U = \sum_{i=1}^n \frac{C_i}{T_i} \leq n(2^{1/n} - 1)$$

n : number of task

C_i : Computation time or Execution time (WCET)

T_i : Release period (P_i)

- If n tends towards infinity

$$\lim_{n \rightarrow \infty} n(2^{1/n} - 1) = \ln 2 \approx 0.693147 \dots$$

- RM Scheduling can meet all of the deadlines if CPU utilization is less than 69.32%.
- 30.7% of the CPU can be used to lower-priority non-real-time tasks.

Shortest Remaining Time First (SRTF) scheduling

- Preemptive version of Shortest Job First
- **Task with the smallest amount of time remaining until completion is selected to execute**
- Decisions are made when
 - Task finished
 - New task added
- Advantages
 - Short tasks are handled very quickly
 - Little overhead through decision steps

Earliest Deadline First (EDF) scheduling

- Dynamic priority scheduling algorithm and preemptive scheduling
- Search for the task closest to its deadline whenever a scheduling event occurs** (task finishes, new task released ...)
- Schedulability
 - If the deadlines equal to their periods, EDF scheduling has a utilization bound of 100%

$$U = \sum_{i=1}^n \frac{C_i}{T_i} \leq 1,$$

n : number of task
 C_i : Computation times or Execution times
 T_i : Release periods (equal to relative deadlines)

- EDF scheduling can guarantee all the deadlines in the system at higher loading !
- Difficult to implement because the relative deadline is not precise to compute
- EDF scheduling is not commonly use in industrial real-time computer systems

	Execution Time	Period
Task 1	1	4
Task 2	1 2	8
Task 3	1 2 3 4 5	16

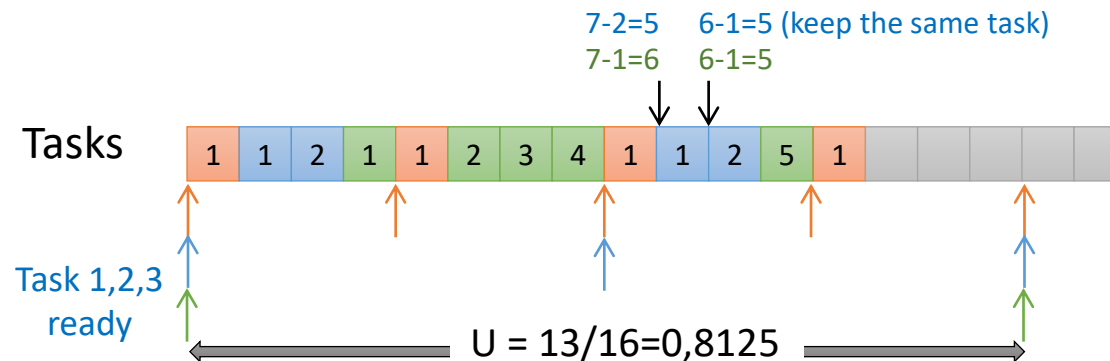
Schedulability :

$$U = (1/4 + 2/8 + 5/16)$$

$$U = (4/16 + 4/16 + 5/16)$$

$$U = 0,8125 < 1$$

Least common multiple of the period : 16



Time Sharing - Round Robin Scheduling

- CPU scheduling algorithm and preemptive scheduling
- **Each task is assigned a fixed time slot in a cyclic way**
- Time slot or Quantum
 - Short quantum : overhead of context switching
 - Long quantum : long response time (infinite quantum = FIFO algorithm)
 - Set a quantum (statistic behavior) when 80% of tasks finish their CPU usage before the end of the quantum

Arrival time = 0

Order : Task 1, Task 2, Task 3

Time slot = 2
(or quantum)




Execution Time

Task 1 1 2 3

Task 2 1 2 3 4 5 6

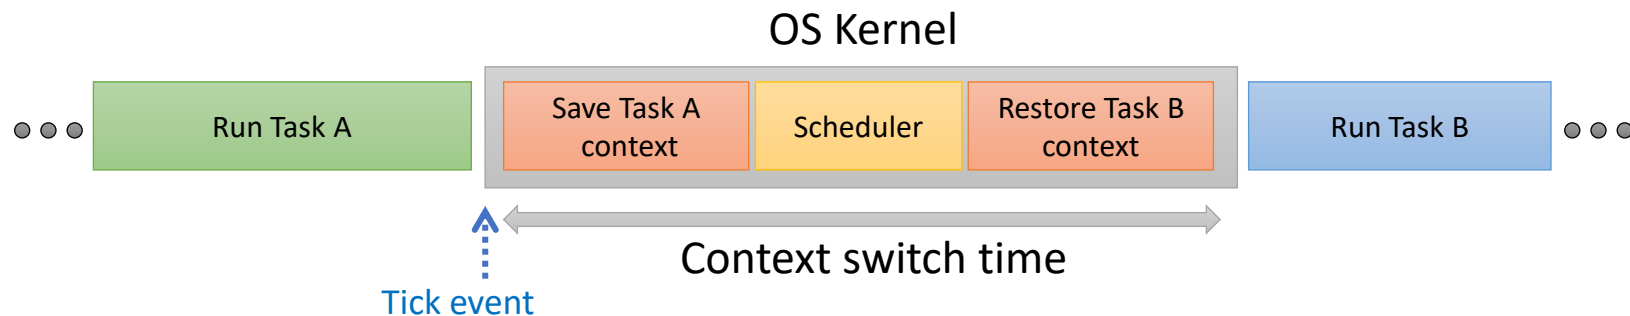
Task 3 1 2 3 4



Context Switch

Context Switch

- A task must be switch out of the CPU to perform another tasks
- Context switch time depends on the architecture
- Save and restore the context
 - On Stack or/and Task Control Block (TCB)
 - Some processors can internally backup and restore the process context

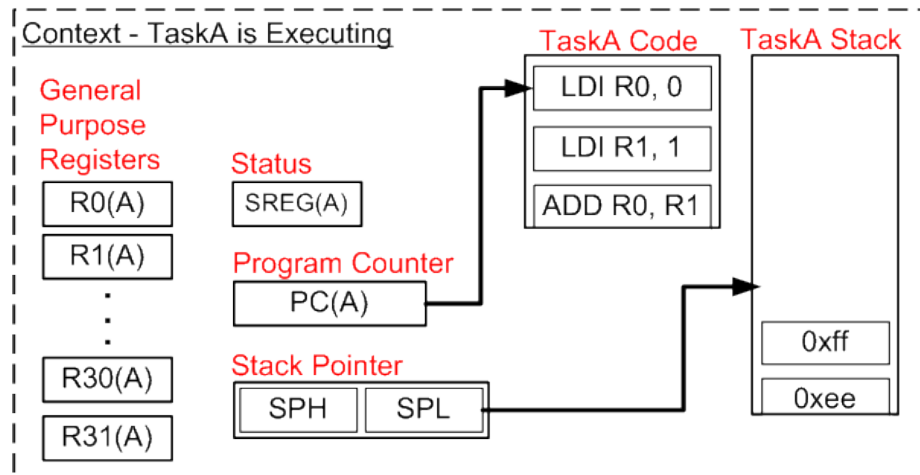


Context Switch – FreeRTOS Example

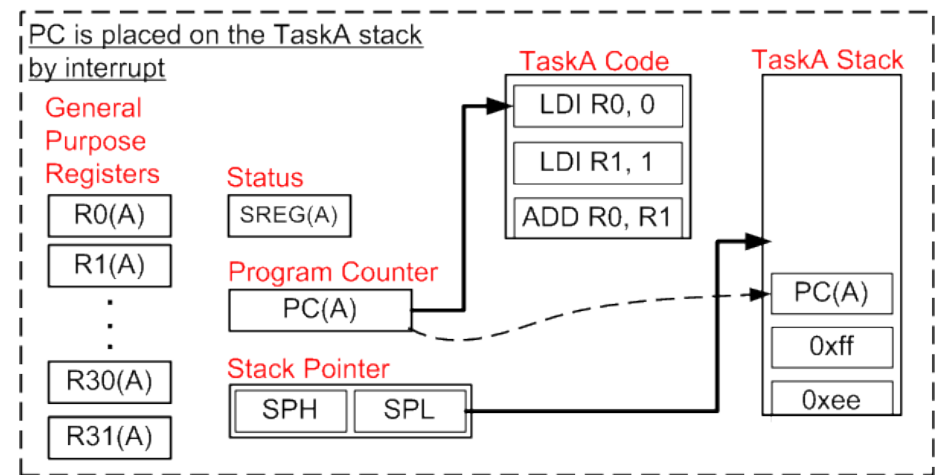
Steps 1 & 2

- Task A is running
- Assume that Task B has previously been suspended (context has already been stored on the Task B stack)
- When the interrupt occurs, the CPU automatically places the current program counter onto the stack before jumping to the start of the RTOS tick ISR.

Step 1



Step 2



Context Switch – FreeRTOS Example

Step 3

- IG_OUTPUT_COMPARE1A() and vPortYieldFromTick() are naked functions
- portSAVE_CONTEXT() pushes the entire CPU execution context onto the stack of Task A and store a copy of the stack pointer

```

/* Interrupt service routine for the RTOS tick. */
void SIG_OUTPUT_COMPARE1A(void) {
    vPortYieldFromTick();
    asm volatile ("reti");
}

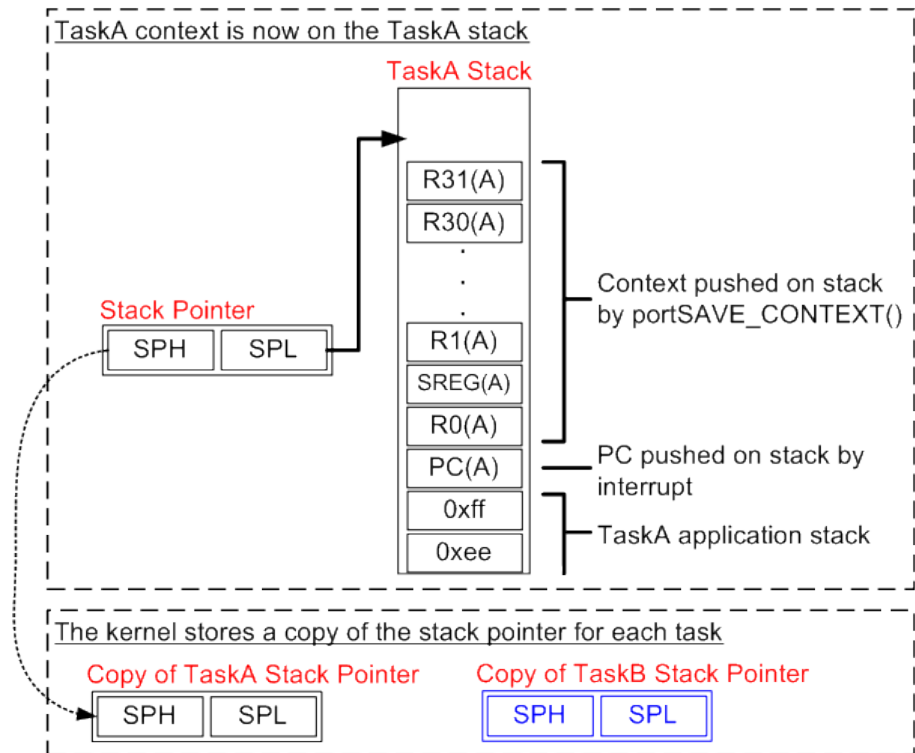
void vPortYieldFromTick(void) {
    portSAVE_CONTEXT();

    vTaskIncrementTick();
    vTaskSwitchContext();
    portRESTORE_CONTEXT();

    asm volatile ("ret");
}

```

Step 3



Context Switch – FreeRTOS Example

Step 4

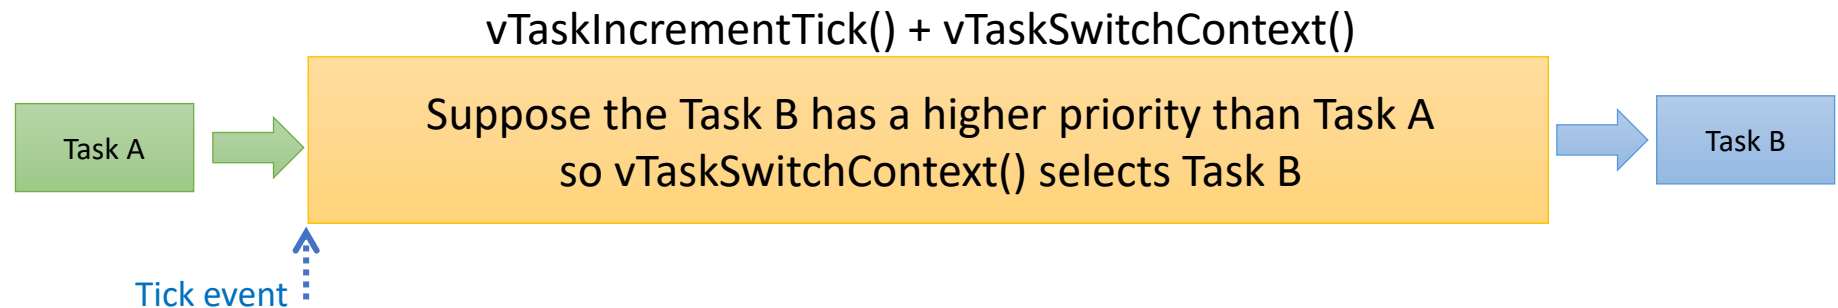
- `vTaskIncrementTick()`
 - Takes care of all aspects related to tick timer
 - Updating the current time
 - Checking whether some task timeouts have expired ...
- `vTaskSwitchContext()`
 - Looks at which tasks are in the ready state
 - Selects the higher priority tasks depending on scheduling policy

```
void vPortYieldFromTick(void) {
    portSAVE_CONTEXT();

    → vTaskIncrementTick();
    → vTaskSwitchContext();
    portRESTORE_CONTEXT();

    asm volatile ("ret");
}
```

Step 4



Context Switch – FreeRTOS Example

Steps 5 & 6

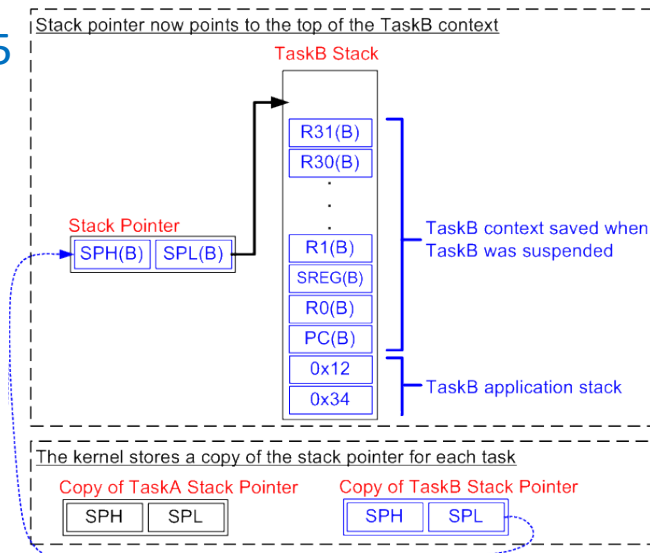
- The Task B context must be restored
- Retrieve the Task B stack pointer from the copy taken when Task B was suspended
- Task B stack pointer is loaded into the processor stack pointer
- CPU stack points to the top of the Task B context
- Restoring the Task B context from its stack into the appropriate processor registers (only the program counter remains on the stack.).

```
void vPortYieldFromTick(void) {
    portSAVE_CONTEXT();

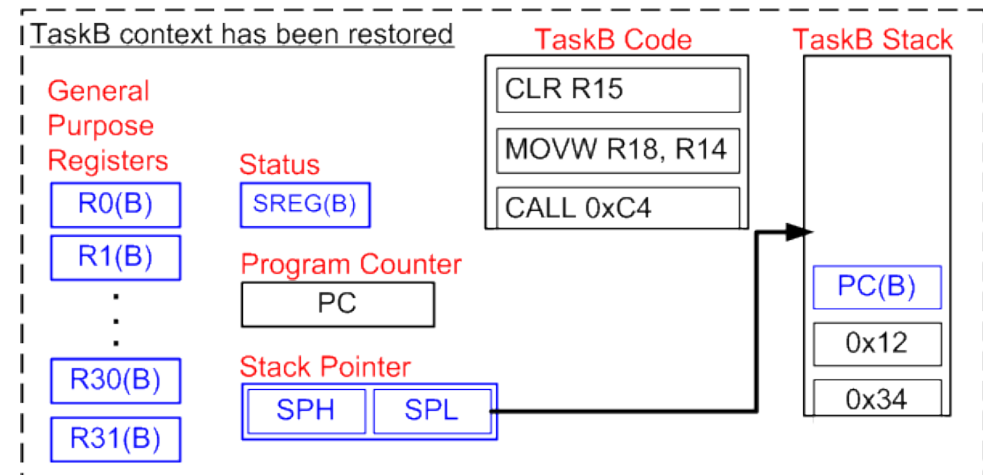
    vTaskIncrementTick();
    vTaskSwitchContext();
    portRESTORE_CONTEXT();

    asm volatile ("ret");
}
```

Step 5



Step 6

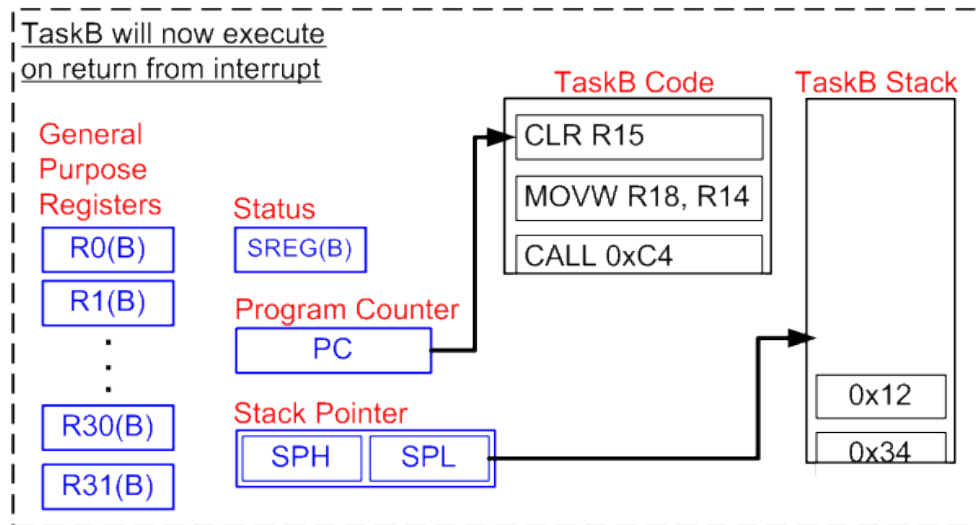


Context Switch – FreeRTOS Example

Step 7

- vPortYieldFromTick() returns to SIG_OUTPUT_COMPARE1A()
- *RETI* instruction assumes the next value on the stack is a return address placed onto the stack when the interrupt occurred
- Task B is running !

Step 7



```
/* Interrupt service routine for the RTOS tick. */
void SIG_OUTPUT_COMPARE1A(void) {
    vPortYieldFromTick();
    asm volatile ("reti");
}

void vPortYieldFromTick(void) {
    portSAVE_CONTEXT();

    vTaskIncrementTick();
    vTaskSwitchContext();
    portRESTORE_CONTEXT();

    asm volatile ("ret");
}
```

Critical Section & Shared Resources

What is a Critical Section ?

- Also called *critical region*
- Useful for concurrent programming
- Protected accesses to a shared resource
 - Data structures
 - Peripheral devices
 - Network connections
- Critical section may be protected by different mechanisms
 - Semaphore
 - Mutex (mutual exclusion)
 - Dedicated functions generally by interrupt masking

Critical Section – P()/V() operations

- Operate on a mutex-semaphore variable
- P operation
 - Dutch word **P**roberen (to attempt)
 - For *wait operation*
 - Require a resource and if not available waits for it
- V operation
 - Dutch word **V**erhogen (to increase)
 - For *signal passing operation*
 - Pass to the OS that the resource is now free to other tasks
- How implemented ?
 - Test & Set instruction
 - Fetch & Add instruction
- Use of standard POSIX 1003.1b, IEEE standard

Critical Section – P()/V() operations

Test & Set instruction

- Implement read/write/test on a shared variable
- Used for Boolean semaphore
- Test & Set instruction
 - Used to both test and (conditionally) write to a memory location as part of a **single atomic (i.e. non-interruptible) operation**
- No other process may begin another test-and-set until the first process's test-and-set is finished
- CPU itself may offer a test-and-set instruction
- <http://en.wikipedia.org/wiki/Test-and-set>

Example

```
boolean lock = false;

void P(void) {
    /* Active wait */
    while (TestAndSet(lock));
}

void V(void) {
    lock = false;
}
```

Critical Section – P()/V() operations

Fetch & Add instruction

- Implement read/write/test on a shared resource
- Used for mutex and counter semaphore
- Fetch & Set instruction
 - atomically (i.e. non-interruptible) modifies the contents of a memory location by a specified value
 - increment the value at address ADDR by VALUE and return the original VALUE at ADDR
- When this instruction is executed by one process in a concurrent system, no other process will ever see an intermediate result
- CPU itself may offer a fetch & Set instruction
- <http://en.wikipedia.org/wiki/Fetch-and-add>

```
int FetchAndAdd(int* addr) {
    int value = *addr;
    *addr = value + 1;
    return value;
}

struct PVType {
    int number;
    int turn;
}
```

Let suppose that
it is an atomic operation

Example of Mutex

```
void PVInit(PVType* pv) {
    pv->number = 0;
    pv->turn = 0;
}


void P(PVType* pv) {
    int turn = FetchAndAdd(&pv.number);
    /* Active wait */
    while (pv.turn != turn);
}

void V((PVType* pv) {
    FetchAndAdd(&pv.turn);
}
```

Shared Resources - Peterson's algorithm

- Formulated by Gary L. Peterson in 1981
- Concurrent programming algorithm for mutual exclusion
- Based on active wait (polling mode)
- Share a single-use resource without conflict from two or more processes

Example for 2 process




Shared Resources - Priority Inversion

- **A high priority task is indirectly preempted by a lower priority task**
- Inverting the relative priorities of the two tasks !
- Does not take care of priority inversion can have disastrous effects
 - Response to emergency situations
 - System can be blocked ...
- Limitation of priority inversion
 - Allow access to critical sections only to tasks with the same priority
 - Use specific semaphores : priority inheritance semaphores or ceiling priority semaphores.

Shared Resources - Priority Inversion

Two tasks, no choice !


- Task 1 acquires the mutex M1
 - An event wakes Task 3 which preempts Task 1
 - Task 3 tries to get the mutex; since it is already acquired by Task 1, Task 3 is suspended.
 - Task 1 runs and releases the mutex 1
 - Task 3 preempts Task 1
- Task 1 was performed first instead of task 3 because of the shared resource !




Shared Resources - Priority Inversion

Three tasks, problem !

- Task 1 acquires the mutex M1
- An event wakes Task 2 which preempts Task 1
- Task 3 becomes in ready stage and preempts Task 2
- Task 3 tries to get the mutex. Since it is already acquired by task 1, Task 3 is suspended and Task 2 runs till ended
- Task 1 runs and releases the mutex 1
- Task 3 can run now till ended

$P(S_x)$: Start of the critical section of mutex x $V(S_x)$: End of the critical section of mutex x  Resource used by task i



The highest priority Task 3 runs last !

Shared Resources - Priority Inversion

Three tasks, Basic Priority inheritance

- Priority inheritance
 - Priority inheritance raises the priority of the blocking task to the blocked one
 - Once the semaphore is released, the blocked task returns to its original priority

$P(S_x)$: Start of the critical section of mutex x

$V(S_x)$: End of the critical section of mutex x




Resource used by task i

Priority(Task 3) > Priority (Task 2) > Priority (Task 1)

Priority(Task 1) = Priority(Task 3)

Original priority of Task 1



The highest priority Task 3 is finished before Task 2



Shared Resources - Priority Inversion

Basic Priority inheritance

- Advantage
 - Bounded Priority inversion
 - Reasonable Run-time performance
- Disadvantage
 - Potential deadlocks
 - Chain-blocking – many preemptions

Shared Resources - Priority Inversion


OCPP - Original Ceiling Priority Protocol

- **A Task 1's priority is raised when a higher-priority Task 2 tries to acquire a resource that Task 1 has locked**
- The task's priority is then raised to the priority ceiling of the resource, ensuring that Task 1 quickly finishes its critical section, unlocking the resource
- A task is only allowed to lock a resource if its dynamic priority is higher than the priority ceilings of all resources locked by other tasks. Otherwise, the task becomes blocked, waiting for the resource
- OCPP changes priority only if an actual block has occurred

Shared Resources - Priority Inversion


Original Ceiling Priority Protocol

Priority(Task 3) > Priority(Task 2) > Priority(Task 1)

 Resource used by task i


Basic Priority inheritance

Task 3 blocked 6 ticks



Original Ceiling Priority inheritance

Task 3 blocked 1 ticks



Shared Resources - Priority Inversion



Immediate Ceiling Priority Protocol

- Also called Highest Locker's priority Protocol (HLP)
- **A task's priority is immediately raised when it locks a resource**
- The task's priority is set to the priority ceiling of the resource, thus no task that may lock the resource is able to get scheduled.
- A task can only lock a resource if its dynamic priority is higher than the priority ceilings of all resources locked by other tasks
- ICPP changes immediately priority

Shared Resources - Priority Inversion


Immediate Ceiling Priority Protocol

Priority(Task 3) > Priority (Task 2) > Priority (Task 1)

 Resource used by task i
 Task i blocked

Original Ceiling Priority inheritance

Task 3 blocked 1 ticks



Immediate Ceiling Priority inheritance

Task 3 blocked 0 ticks

